

Basic Programming

Lesson 10. Analyzing dataframes. Pandas and Matplotlib Python libraries

Alexandra Ciobotaru

Syllabus

- **Lesson 1. Computers, Programming and Cognitive Science. From pseudocode to programming languages.**
- **Lesson 2. Variables in Python. Basic calculus. Using Math library, type() and help() functions.**
- **Lesson 3. Collections: Lists, Tuples.**
- **Lesson 4. Strings. Working with strings.**
- **Lesson 5. Branching and decisions: Logical operators, If-Statements, Nested conditions. Loops: For and While**
- **Lesson 6. Lesson 5 continued Quiz 1 (25%).**
- **Lesson 7. Creating functions. Recursive functions. Matrices.**
- **Lesson 8. Collections: Dictionaries. JSON construction.**
- **Lesson 9. Working with files. Reading and Writing.**
- **Lesson 10. Analyzing dataframes. Pandas and Matplotlib Python libraries.**
- **Lesson 12. Object-Oriented Programming: Encapsulation, Inheritance and Polymorphism. Quiz 2 (25%).**
- **Lesson 13. Object-Oriented Programming (cont.). Error handling. Best practices when programming.**
- **Lab (20%) + final exam (30%).**

What you'll learn today:

- 1. What is Pandas
 - 1.1. Pandas Series
 - 1.2. Creating a dataframe from series
 - 1.3. Add series to dataframe
- 2. Dataframes
 - 2.1. Creating dataframes
 - 2.2. Loading csv data into dataframes
 - 2.3. Selecting specific elements from dataframe
 - 2.4. Dataframe statistics
 - 2.5. Writing dataframe to csv
 - 2.6. Read and write to excel
- 3. Matplotlib library
 - 3.1. Bar plots
 - 3.2. Line plots
 - 3.3. Pie chart

But first

...A small remaining from last courses: Python **enumerate()**

```
>>> for count, value in enumerate(values):  
...     print(count, value)  
...  
0 a  
1 b  
2 c
```

When you use `enumerate()`, the function gives you back two loop variables:

- The **count** of the current iteration
- The **value** of the item at the current iteration

Of course, as always, you can name your variables how you wish.

What is Pandas?



- A library for manipulating spreadsheets (tabular data)
- Useful for data preparation, data cleaning and analysis
- Installation: `pip install pandas`
- Importing: `import pandas as pd`

Pandas Series

- One dimensional datastructures
- A pandas series returns an object in the form of a list, having index starting from 0 to n
- Creating a Series from a list (by default, indexing is done from 0)

```
list_1 = ['a', 'b', 'c', 'd']  
ser = pd.Series(list_1)  
print(ser)
```

The diagram illustrates the structure of a Pandas Series. It features a central column of names, each preceded by an index value from 0 to 9. The names are: Lois, Brenda, Joe, Diane, Benjamin, Patrick, Nancy, Carol, Frances, and Diana. The index values are highlighted in light blue boxes. A large black bracket on the left side of the index column is labeled 'Index' in red text. A large black bracket on the right side of the name column is labeled 'Data' in red text. The entire structure is enclosed in a black rectangular frame.

	First Name
0	Lois
1	Brenda
2	Joe
3	Diane
4	Benjamin
5	Patrick
6	Nancy
7	Carol
8	Frances
9	Diana

- Create Series with custom index

```
labels = [1,2,3,4]  
ser_1 = pd.Series(data=list_1, index=labels)
```

- Create series from numpy array

```
arr_1 = np.array([1,2,3,4])  
ser_2 = pd.Series(arr_1)
```

- Create series from dictionary

```
dict_1 = {'f_name':'Derek', 'l_name':'Barnabas',  
          'age':44}  
ser_3 = pd.Series(dict_1)
```

- Create series from scalar

```
s = pd.Series(7, index=[0, 1, 2, 3])
```

Creating a dataframe from multiple series

Series 1		Series 2		Series 3		DataFrame			
Mango		Apple		Banana		Mango	Apple	Banana	
0	4	0	5	0	2	0	4	5	2
1	5	1	4	1	3	1	5	4	3
2	6	2	3	2	5	2	6	3	5
3	3	3	0	3	2	3	3	0	2
4	1	4	2	4	7	4	1	2	7

Image source: <https://www.datasciencemadesimple.com/create-series-in-python-pandas/>

Creating a dataframe from multiple series

```
mango = [4, 5, 6, 3, 1]
apple = [4, 4, 3, 0, 2]
banana = [2, 3, 5, 2, 7]
mango_series = pd.Series(mango)
apple_series = pd.Series(apple)
banana_series = pd.Series(banana)
frame = {'Mango': mango_series, 'Apple': apple_series,
        'Banana': banana_series}
result = pd.DataFrame(frame)
print(result)
```

Add series to dataframe (add a column)

```
rotten = [0, 1, 0, 1, 0]  
result['Rotten'] = rotten  
print(result)
```

	Mango	Apple	Banana
0	4	4	2
1	5	4	3
2	6	3	5
3	3	0	2
4	1	2	7



	Mango	Apple	Banana	Rotten
0	4	4	2	0
1	5	4	3	1
2	6	3	5	0
3	3	0	2	1
4	1	2	7	0

Dataframes

- **Pandas DataFrame** is two-dimensional size, mutable, potentially heterogeneous tabular data structure with labeled axes (rows and columns).
- Pandas DataFrame consists of three principal components, the **data**, **rows**, and **columns**.
- We can perform many operations on these datasets like arithmetic operation, columns/rows selection, columns/rows addition etc.

The diagram illustrates a Pandas DataFrame with the following structure:

	Name	Team	Number	Position	Age
0	Avery Bradley	Boston Celtics	0.0	PG	25.0
1	John Holland	Boston Celtics	30.0	SG	27.0
2	Jonas Jerebko	Boston Celtics	8.0	PF	29.0
3	Jordan Mickey	Boston Celtics	NaN	PF	21.0
4	Terry Rozier	Boston Celtics	12.0	PG	22.0
5	Jared Sullinger	Boston Celtics	7.0	C	NaN
6	Evan Turner	Boston Celtics	11.0	SG	27.0

Annotations in the diagram:

- Columns:** A blue arrow points to the column headers (Name, Team, Number, Position, Age).
- Rows:** An orange arrow points to the row indices (0, 1, 2, 3, 4, 5, 6).
- Data:** A pink box highlights the data cells, including the values and the missing values (NaN).
- Missing value:** A blue arrow points to the NaN value in the 'Number' column for row 3.

Logo: æG

Creating a dataframe

- Creating an empty dataframe

```
df = pd.DataFrame()
```

- Creating a dataframe from a dictionary

```
# Let's define a dictionary containing employee data
data = {'Name': ['Jai', 'Princi', 'Gaurav', 'Anuj'],
        'Age': [27, 24, 22, 32],
        'Address': ['Delhi', 'Kanpur', 'Allahabad', 'Kannauj'],
        'Qualification': ['Msc', 'MA', 'MCA', 'Phd']}
# Convert the dictionary into DataFrame
df = pd.DataFrame(data)
print(df)
```

- You can see the first 5 rows of the dataset:

```
print(df.head())
```

- Values in dictionary must have the same length, otherwise the program will throw an error.
- Other methods that show basic info about dataframes:

```
df.shape -> (rows, columns)
```

```
df.index -> describes index
```

```
df.count() -> no of non-NA values
```

Loading data from a csv into a dataframe

- CSV – Comma Separated files
- Loading csv data into a dataframe is done using **.read_csv** method:

```
data = pd.read_csv("nba.csv")
```
- You can also specify the column for indexing:

```
data = pd.read_csv("nba.csv", index_col = "Name")
```
- View all columns:

```
for col in data.columns:  
    print(col)
```
- Delete rows with empty values:

```
data = data.dropna()
```
- **Exercise:** can you define a list 'col' that holds all column names using accumulation? Can you do it also using list comprehension?
- Another way of doing this is using column.values method:

```
print(list(data.columns.values))
```

Selecting specific elements in dataframe

- You can select rows using **.loc** method on the index elements:

```
first = data.loc["Avery Bradley"]  
second = data.loc["R.J. Hunter"]  
print(first)  
print(second)
```

loc -> selects [row_label, column_label] iloc -> selects [row_position, column_position]

- Select element at specific row & column:

```
df = pd.read_csv('nba.csv', index_col='Name')  
print(df.loc['Avery Bradley', 'Weight'])
```

- You can select elements by indexes using **.iloc** method:

```
print(df.iloc[2, 6])
```

- You can also select by condition:

```
new_df = df[df["Team"] == "Boston Celtics"]
```

Selecting specific columns in csv

- We can select columns in dataframe by calling them by their column name, using the [] operator

```
df = pd.read_csv('nba.csv')
```

- We can print the columns just to remember what's inside

```
print(df.columns)
```

- Print the contents of column Name:

```
print(df["Name"])
```

- Selecting more than one column is done using [][] operator

- Print the contents of columns Name and Age:

```
print(df[["Name", "Age"]])
```

- Or, we can select specific columns when loading the csv:

```
df = pd.read_csv('nba.csv')
```

```
df = pd.DataFrame(df, columns=['Name', 'Team'])
```

Selecting with slices

	Weather	Temperature	Wind	Humidity
Day				
Mon	Sunny	12.79	13	30
Tue	Sunny	19.67	28	96
Wed	Sunny	17.51	16	20
Thu	Cloudy	14.44	11	22
Fri	Shower	10.51	26	79
Sat	Shower	11.07	27	62
Sun	Sunny	17.50	20	10

df

```
cols = ['Temperature', 'Wind']  
df.loc['Mon':'Thu', cols]
```

	Temperature	Wind
Day		
Mon	12.79	13
Tue	19.67	28
Wed	17.51	16
Thu	14.44	11

Image source: <https://towardsdatascience.com/how-to-use-loc-and-iloc-for-selecting-data-in-pandas-bd09cb4c3d79>

- Selecting rows with slices:
- `df = pd.read_csv('nba.csv')`
- `new_df = df[:5]` # will select the first 5 rows and put them in new_df
- **Sorting:** `df.sort_values(by='Salary')`

Deleting specific columns or rows

- In Order to delete a column in Pandas DataFrame, we can use the drop() method.

```
df.drop(["Team", "Weight"], axis = 1, inplace=True)
```

- Inplace = True means that the changes are made in the original dataframe; False means it creates a new dataframe
- Axis = 1 means that dropping is done on columns
- Axis = 0 means that dropping is done on rows, by index
- Delete a few specified rows at index values 1, 2, 4.

```
# Note that the index values do not always align to row numbers.
```

```
data = data.drop(labels=[1,2,4], axis=0)
```

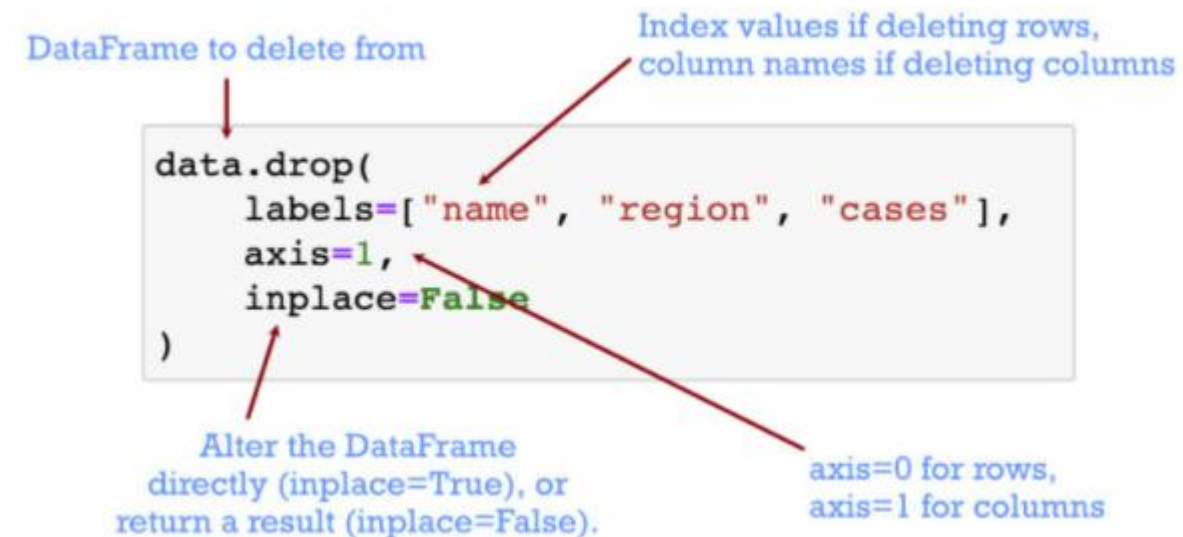


Image source: <https://www.shanelynn.ie/pandas-drop-delete-dataframe-rows-columns/>

Dataframe statistics

- If we have the dataframe:

```
data = pd.read_csv("nba.csv", index_col = "Name")
```

- The following methods compute on all possible columns:

```
data.sum() -> sum of values
```

```
data.min() / data.max() -> min/max values
```

```
data.idxmin() / data.idxmax() -> min/max index value
```

```
data.describe() -> summary statistics
```

```
data.mean() -> computes mean on all possible columns
```

```
data.median() -> computes median on all possible columns
```

Writing dataframe to csv

- After you modified the dataframe, you can write it to a csv using **.to_csv** method:
- # do something to df dataframe and then save it
- `df.to_csv('name_of_csv.csv')`
- Encoding parameter – useful to set it to 'utf-8' when dealing with Romanian, both for reading and for writing csvs.

Read and write to Excel

- You can open an excel file using **.read_excel** method:

```
df = pd.read_excel('file.xlsx')
```

- You can create an excel file from a dataframe using **.to_excel** method:

```
df.to_excel('path_to_excel_file.xlsx', sheet_name =  
    'Sheet1')
```

- You can read each sheet in a different dataframe:

```
df = pd.read_excel(xlsx, 'Sheet1')
```

What is Matplotlib?



- A comprehensive library for creating static, animated, and interactive visualizations in Python
- Installation: `pip install matplotlib`
- Importing: `import matplotlib.pyplot as plt`
- You can create plots with: matplotlib, seaborn and pandas as well!

Bar plot

- You can use **.plot.bar()** method to plot the graph vertically in form of rectangular bars

```
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
```

- Create a dataframe from a random numpy array

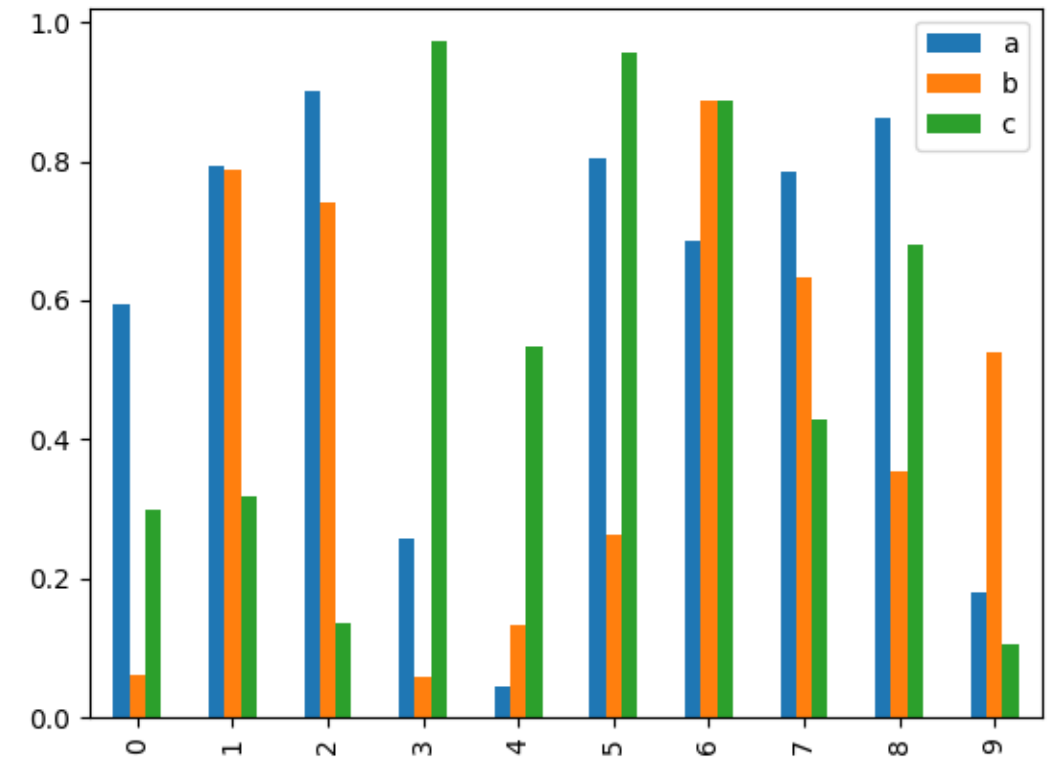
```
rnd_arr = np.random.rand(10, 3)
print(rnd_arr)
df = pd.DataFrame(rnd_arr, columns=['a', 'b', 'c'])
df.plot.bar()
plt.show()
```

- Plotting two columns from dataframe:

```
df.plot.bar(x = 'name_of_column1', 'name_of_column2')
```

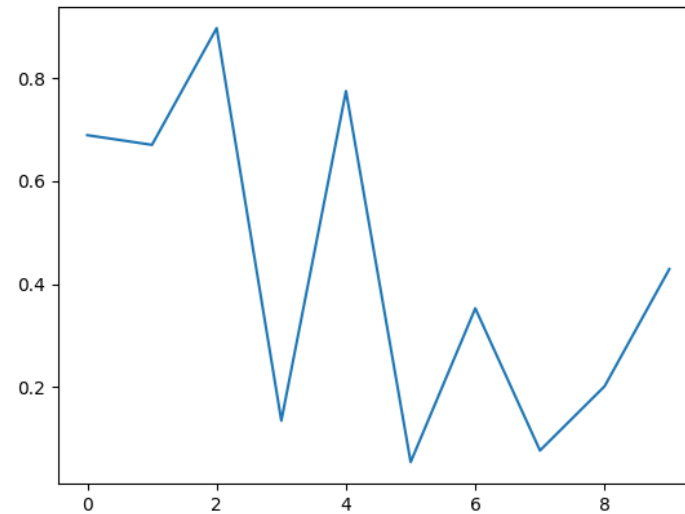
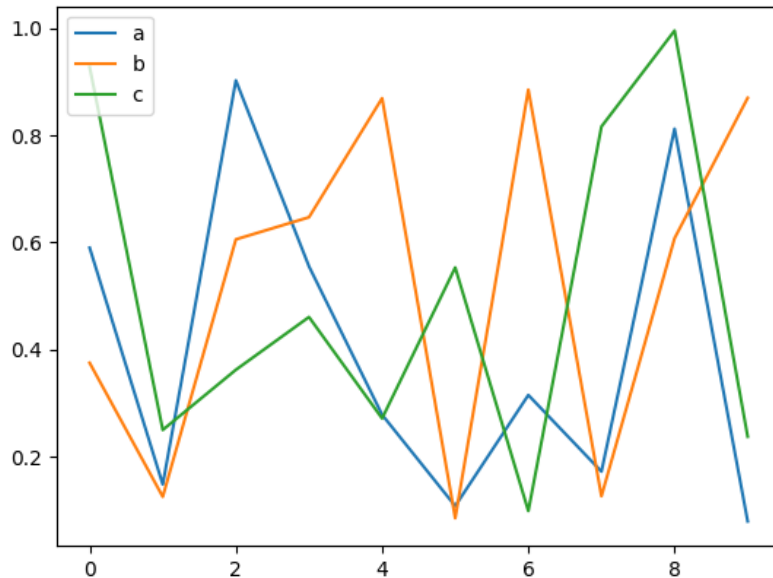
Bar plot

	a	b	c
0	0.594736	0.060920	0.299605
1	0.792587	0.788698	0.317061
2	0.900321	0.740707	0.136934
3	0.257342	0.058175	0.972296
4	0.045137	0.131768	0.534247
5	0.803600	0.262312	0.957628
6	0.684835	0.888481	0.887564
7	0.784256	0.633894	0.429599
8	0.862750	0.354024	0.679965
9	0.181168	0.524034	0.106147



Line plot

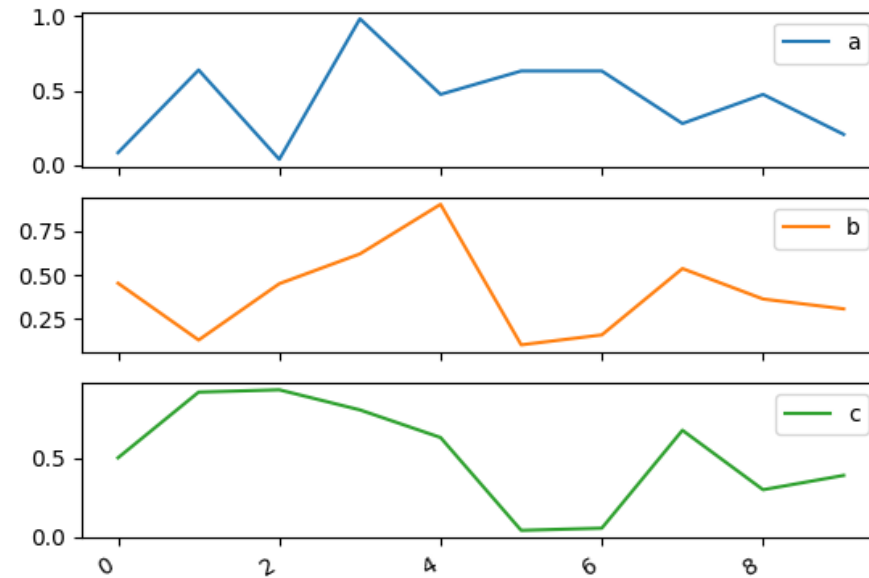
- On the same dataframe as previous, if I write `df.plot.line()` instead of `df.plot.bar()`:
- If I want, I can select only one column to plot:
`df['a'].plot.line()`



Disclaimer: plots are different because of creating another random array on each run

- We can also see each line in dataframe separately:

```
axes = df.plot.line(subplots=True)
```



Creating line plots with matplotlib

```
import matplotlib.pyplot as plt
import numpy as np
```

By setting coordinates for each point:

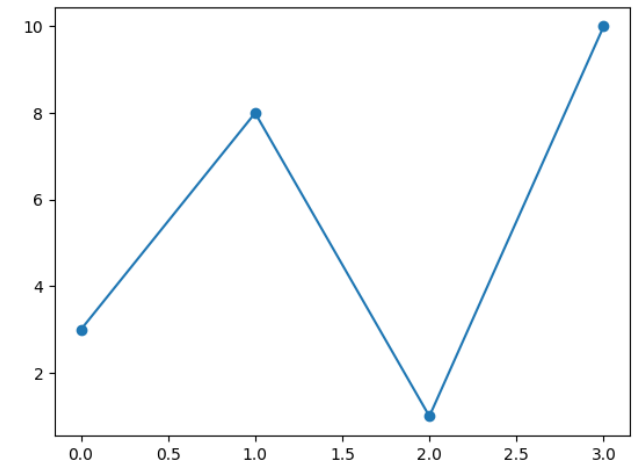
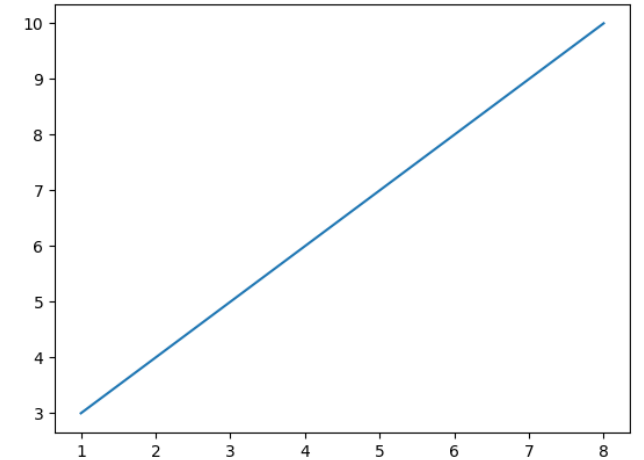
```
xpoints = np.array([1, 8])
ypoints = np.array([3, 10])
plt.plot(xpoints, ypoints)
plt.show()
```

By setting only y coordinates.

And you can also use markers.

```
ypoints = np.array([3, 8, 1, 10])
plt.plot(ypoints, marker = 'o')
plt.show()
```

More on matplotlib: https://www.w3schools.com/python/matplotlib_intro.asp



Pie chart

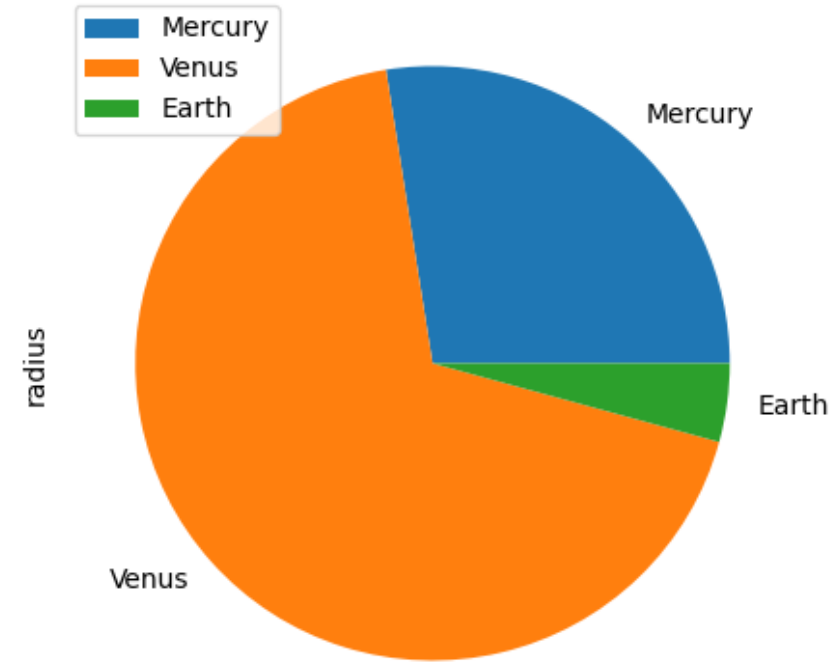
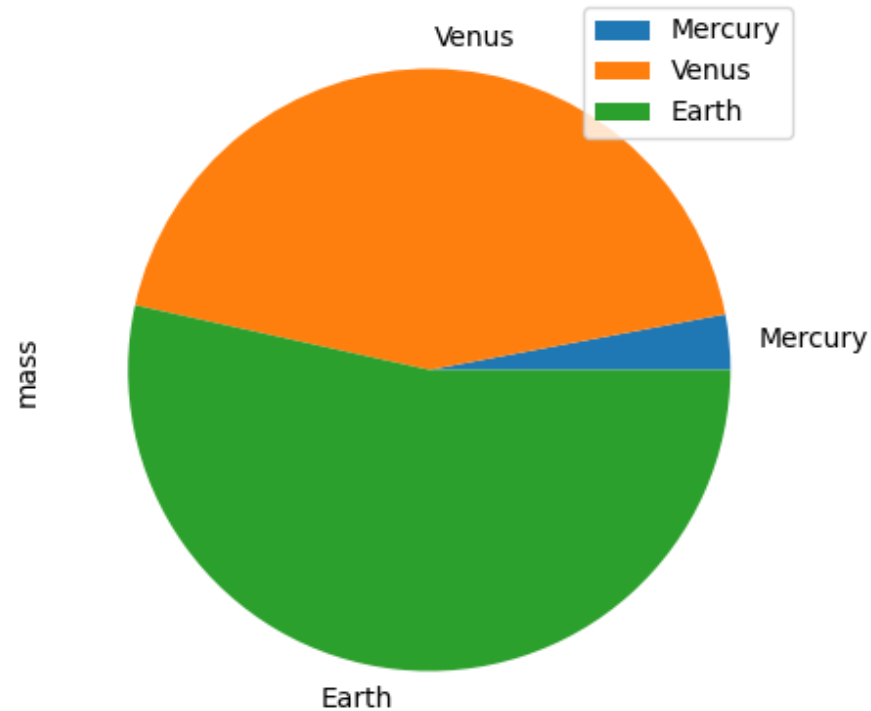
- Here is a dataframe holding information about some planets mass and radius:

```
df = pd.DataFrame({'mass': [0.330, 4.87, 5.97],  
                   'radius': [2439.7, 6051.8, 378.1]},  
                  index=['Mercury', 'Venus', 'Earth'])
```

- You can use `.plot.pie` method to get a pie plot:

```
plot = df.plot.pie(y='mass', figsize=(5, 5))
```

Pie chart



More about pandas plots: https://pandas.pydata.org/pandas-docs/stable/user_guide/visualization.html

Thank you!